



# SAIR

Spatial AI & Robotics Lab

# CSE 473/573-A

## L11: ALIGNMENT & FITTING

Chen Wang

Spatial AI & Robotics Lab

Department of Computer Science and Engineering




**University at Buffalo** The State University of New York

# What are Alignment and Fitting?

---


- Alignment
  - Find the parameters of a transformation that best aligns matched points
- Fitting
  - Find the parameters of a model that best fit the data









# Fitting and Alignment: Methods

---

- Hypothesize and test
  - Generalized Hough transform
- **General Alignment**
  - **Homographies**
  - **Rotational Panoramas**
  - Global Alignment
  - RANSAC
  - Warping
  - Blending
- Global optimization / Search for parameters
  - Least squares fit
  - Robust least squares
  - Other parameter search methods



# Motivation: Medical image registration



# Motivation: Mosaics

---

- Getting the whole picture
  - Typical camera:  $50^\circ \times 35^\circ$



# Motivation: Mosaics

---

- Getting the whole picture
  - Typical camera:  $50^\circ \times 35^\circ$
  - Human Vision:  $176^\circ \times 135^\circ$






# Motivation: Mosaics

---

- Getting the whole picture
  - Typical camera:  $50^\circ \times 35^\circ$
  - Human Vision:  $176^\circ \times 135^\circ$



# Alignment

- Homography
- Rotational Panoramas
- RANSAC (Next Lecture)
- Global alignment
- Warping
- Blending



(a)



(b)






# Motion models

- What happens when we take two images with a camera and try to align them?
- translation?
- rotation?
- scale?
- affine?
- perspective?




# Image Warping (Recap)

- image filtering: change *range* of image
  - $g(x) = h(f(x))$




- image warping: change *domain* of image
  - $g(x) = f(h(x))$




# Image Warping

- image filtering: change *range* of image
  - $g(x) = h(f(x))$



- image warping: change *domain* of image
  - $g(x) = f(h(x))$



# Parametric (global) warping

- Examples of parametric warps:



translation



rotation



aspect



affine




perspective



cylindrical


# Image Warping

- Given a coordinate transform  $\mathbf{x}' = \mathbf{h}(\mathbf{x})$  and a source image  $\mathbf{f}(\mathbf{x})$ , how do we compute a transformed image  $\mathbf{g}(\mathbf{x}') = \mathbf{f}(\mathbf{h}(\mathbf{x}))$ ?



# Forward Warping


- Send each pixel  $f(\mathbf{x})$  to its corresponding location  $\mathbf{x}' = h(\mathbf{x})$  in  $g(\mathbf{x}')$
- What if pixel lands “between” two pixels?






# Forward Warping

- Send each pixel  $f(\mathbf{x})$  to its corresponding location  $\mathbf{x}' = h(\mathbf{x})$  in  $g(\mathbf{x}')$
- What if pixel lands “between” two pixels?
- Answer: add “contribution” to several pixels, normalize later






# Interpolation

- Possible interpolation filters:
  - **Nearest Neighbor**
  - **Bilinear**
  - **Bicubic**
- Needed to prevent “jaggies” and “texture crawl”




# Bilinear interpolation



# Bilinear interpolation



Sampling at  $f(x,y)$ :



$$\begin{aligned} f(x, y) = & (1 - a)(1 - b) f[i, j] \\ & + a(1 - b) f[i + 1, j] \\ & + ab f[i + 1, j + 1] \\ & + (1 - a)b f[i, j + 1] \end{aligned}$$

# Bicubic interpolation

- Bilinear interpolation processes 2x2 (4 pixels) squares
- Bicubic interpolation processes 4x4 (16 pixels) squares to take image gradients into account.



# Bicubic interpolation

$$p(x, y) = \sum_{i=0}^3 \sum_{j=0}^3 a_{ij} x^i y^j.$$

surface  $p(x, y)$  on the **unit square**  $[0, 1] \times [0, 1]$  that is continuous

This requires determining the 16 coefficients.

Consider 4 corners of the unit square.  $(0, 0)$   $(1, 0)$   $(0, 1)$   $(1, 1)$

1.  $f(0, 0) = p(0, 0) = a_{00},$
2.  $f(1, 0) = p(1, 0) = a_{00} + a_{10} + a_{20} + a_{30},$
3.  $f(0, 1) = p(0, 1) = a_{00} + a_{01} + a_{02} + a_{03},$
4.  $f(1, 1) = p(1, 1) = \sum_{i=0}^3 \sum_{j=0}^3 a_{ij}.$



# Bicubic interpolation

We need following derivatives

$$p_x(x, y) = \sum_{i=1}^3 \sum_{j=0}^3 a_{ij} i x^{i-1} y^j,$$

$$p_y(x, y) = \sum_{i=0}^3 \sum_{j=1}^3 a_{ij} x^i j y^{j-1},$$

$$p_{xy}(x, y) = \sum_{i=1}^3 \sum_{j=1}^3 a_{ij} i x^{i-1} j y^{j-1}.$$

Likewise, eight equations for the derivatives in the  $x$  and the  $y$  directions:

1.  $f_x(0, 0) = p_x(0, 0) = a_{10},$
2.  $f_x(1, 0) = p_x(1, 0) = a_{10} + 2a_{20} + 3a_{30},$
3.  $f_x(0, 1) = p_x(0, 1) = a_{10} + a_{11} + a_{12} + a_{13},$
4.  $f_x(1, 1) = p_x(1, 1) = \sum_{i=1}^3 \sum_{j=0}^3 a_{ij} i,$
5.  $f_y(0, 0) = p_y(0, 0) = a_{01},$
6.  $f_y(1, 0) = p_y(1, 0) = a_{01} + a_{11} + a_{21} + a_{31},$
7.  $f_y(0, 1) = p_y(0, 1) = a_{01} + 2a_{02} + 3a_{03},$
8.  $f_y(1, 1) = p_y(1, 1) = \sum_{i=0}^3 \sum_{j=1}^3 a_{ij} j.$

And four equations for the  $xy$  mixed partial derivative:

1.  $f_{xy}(0, 0) = p_{xy}(0, 0) = a_{11},$
2.  $f_{xy}(1, 0) = p_{xy}(1, 0) = a_{11} + 2a_{21} + 3a_{31},$
3.  $f_{xy}(0, 1) = p_{xy}(0, 1) = a_{11} + 2a_{12} + 3a_{13},$
4.  $f_{xy}(1, 1) = p_{xy}(1, 1) = \sum_{i=1}^3 \sum_{j=1}^3 a_{ij} ij.$

# Bicubic interpolation

Grouping the unknown parameters  $a_{ij}$  in a vector

$$\alpha = [a_{00} \ a_{10} \ a_{20} \ a_{30} \ a_{01} \ a_{11} \ a_{21} \ a_{31} \ a_{02} \ a_{12} \ a_{22} \ a_{32} \ a_{03} \ a_{13} \ a_{23} \ a_{33}]^T$$

and letting


$$x = [f(0,0) \ f(1,0) \ f(0,1) \ f(1,1) \ f_x(0,0) \ f_x(1,0) \ f_x(0,1) \ f_x(1,1) \ f_y(0,0) \ f_y(1,0) \ f_y(0,1) \ f_y(1,1) \ f_{xy}(0,0) \ f_{xy}(1,0) \ f_{xy}(0,1) \ f_{xy}(1,1)]^T,$$

the above system of equations can be reformulated into a matrix for the linear equation  $A\alpha = x$ .

Inverting the matrix gives the more useful linear equation  $A^{-1}x = \alpha$ , where

$$A^{-1} = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ -3 & 3 & 0 & 0 & -2 & -1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2 & -2 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & -3 & 3 & 0 & 0 & -2 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 2 & -2 & 0 & 0 & 1 & 1 & 0 & 0 \\ -3 & 0 & 3 & 0 & 0 & 0 & 0 & 0 & -2 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & -3 & 0 & 3 & 0 & 0 & 0 & 0 & 0 & -2 & 0 & -1 & 0 \\ 9 & -9 & -9 & 9 & 6 & 3 & -6 & -3 & 6 & -6 & 3 & -3 & 4 & 2 & 2 & 1 \\ -6 & 6 & 6 & -6 & -3 & -3 & 3 & 3 & -4 & 4 & -2 & 2 & -2 & -2 & -1 & -1 \\ 2 & 0 & -2 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 2 & 0 & -2 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \\ -6 & 6 & 6 & -6 & -4 & -2 & 4 & 2 & -3 & 3 & -3 & 3 & -2 & -1 & -2 & -1 \\ 4 & -4 & -4 & 4 & 2 & 2 & -2 & -2 & 2 & -2 & 2 & -2 & 1 & 1 & 1 & 1 \end{bmatrix},$$

which allows  $\alpha$  to be calculated quickly and easily.



# Panoramas




image from S. Seitz

Obtain a wide angle view by combining multiple images.

# How to stitch together a panorama?

---

- Basic Procedure

- Take a sequence of images from the same position
    - Rotate the camera about its **optical center**.
  - Compute transformation between 2<sup>nd</sup> and 1<sup>st</sup> image
  - Transform the 2<sup>nd</sup> image to overlap with the 1<sup>st</sup>
  - Blend the two together to create a mosaic
  - (If there are more images, repeat)
- Why should this work at all?
    - What about the 3D geometry of the scene?
    - Why aren't we using it?

# Correspondence

- Allows us to map image back to some real space






image from S. Seitz


# Panoramas: generating synthetic views



Can generate any synthetic camera view  
as long as it has **the same center of projection!**



# Image reprojection



- The mosaic has a natural interpretation in 3D
  - The images are reprojected onto a common plane
  - The mosaic is formed on this plane
  - Mosaic is a *synthetic wide-angle camera*

# Recap: 2D coordinate transformations

- translation:  $\mathbf{x}' = \mathbf{x} + \mathbf{t}$   $\mathbf{x} = (x, y)$
- rotation:  $\mathbf{x}' = \mathbf{R} \mathbf{x} + \mathbf{t}$
- similarity:  $\mathbf{x}' = s \mathbf{R} \mathbf{x} + \mathbf{t}$
- affine:  $\mathbf{x}' = \mathbf{A} \mathbf{x} + \mathbf{t}$
- perspective:  $\underline{\mathbf{x}}' \cong \mathbf{H} \underline{\mathbf{x}}$   $\underline{\mathbf{x}} = (x, y, 1)$   
( $\underline{\mathbf{x}}$  is a *homogeneous* coordinate)
- These all form a nested *group* (closed w/ inv.)

# Basic 2D Transformations

- Basic 2D transformations as 3x3 matrices

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$


Rotate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear (Skew)

# Image alignment

- Two broad approaches:
  - Direct (pixel-based) alignment
    - Search for alignment where most pixels agree
  - Feature-based alignment
    - Search for alignment where *extracted features* agree
    - Can be verified using pixel-based alignment




# Fitting an affine transformation



Affine model approximates perspective projection of planar objects.

# Fitting an affine transformation

- Assuming we know the correspondences, how do we get the transformation?




$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$



# Fitting an affine transformation

- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\begin{bmatrix}$$

$$\begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix}$$

# Fitting an affine transformation


$$\begin{bmatrix} \dots & & & & & \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

- How many matches (correspondence pairs) do we need to solve for the transformation parameters?
- Once we have solved for the parameters, how do we compute the coordinates of the corresponding point for  $(x_{new}, y_{new})$ ?

# Homography

- Take as a 2D **image warp** using projective transform.
- A **projective transform** is a mapping between any two PPs with the same center of projection
  - rectangle should map to arbitrary quadrilateral
  - parallel lines aren't preserved.
  - but straight lines are preserved.
- Called **Homography**

$$\begin{bmatrix} wx' \\ wy' \\ w \\ \mathbf{p}' \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ \mathbf{H} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \\ \mathbf{p} \end{bmatrix}$$



# Solving for homographies

$(x, y)$



$$\begin{pmatrix} wx'/w & wy'/w \end{pmatrix} = (x', y')$$


To **apply** a given homography  $\mathbf{H}$

- Compute  $\mathbf{p}' = \mathbf{H}\mathbf{p}$  (regular matrix multiply)
- Convert  $\mathbf{p}'$  from homogeneous to image coordinates

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$\mathbf{p}'$   $\mathbf{H}$   $\mathbf{p}$

# Solving for homographies



To **compute** the homography given pairs of corresponding points, we need to set up an equation where the parameters of **H** are the unknowns...

# Solving for homographies

$$\mathbf{p}' = \mathbf{H}\mathbf{p} \quad \begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Can set scale factor  $i = 1$  or  $\|\mathbf{H}\| = 1$ . So, there are 8 unknowns.
- Set up a system of linear equations:

$$\bullet \mathbf{A}\mathbf{h} = \mathbf{b}$$

where vector of unknowns  $\mathbf{h} = [a, b, c, d, e, f, g, h]^T$

- Need at least 8 equations (4 points), but the more the better...
- Solve for  $\mathbf{H}$ . If over-constrained, solve using least-squares:

$$\min \|\mathbf{A}\mathbf{h} - \mathbf{b}\|^2$$

$$\mathbf{h} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}$$

# Proof of least squares

- $F(h) = ||Ah - b||^2 = (Ah - b)^T (Ah - b)$
- $F(h) = h^T A^T Ah - h^T A^T b - b^T Ah + b^T b$
- $\frac{\partial}{\partial h} F(h) = 2A^T Ah - A^T b - (b^T A)^T$
- Setting derivative to 0:  $\frac{\partial}{\partial h} F(h) = 0$
- $A^T Ah = A^T b$
- $h = (A^T A)^{-1} A^T b$



# How to stitch together a panorama?

---

- Take a sequence of images from the same position
  - Rotate the camera about its optical center
- Compute transformation between second image and first
- Transform the second image to overlap with the first
- Blend the two together to create a mosaic
- If there are more images, repeat

# Content

---

- Stitching
  - Alignment
    - Interpolation
    - Homography
  - Fitting
    - Solving for homographies

