



# SAIR

Spatial AI & Robotics Lab

# CSE 473/573-A

## L16: SEGMENTATION

Chen Wang

Spatial AI & Robotics Lab


Department of Computer Science and Engineering



**University at Buffalo** The State University of New York

# Image Segmentation

- **Partition** the pixels of an image into **groups** that strongly correlate with the objects in an image
- Typically, the first step in any automated computer vision application.



# Why & What?


---

- Useful mid-level representation of an image
  - Can facilitate better downstream tasks
- Segmentation should be homogeneous with some characteristic (gray level, texture, color, motion)
- The desired type segmentation depends on the task
  - Broad theory is absent at present
- Variety of approaches, algorithms, and applications
  - Finding people, summarizing video, annotation figures, background subtraction, finding buildings/rivers in satellite images.

# Thought Exercise....

---


- Based on what you already know about CVIP
  - How would you do segmentation?
- Did you come up with
  - Feature Detection?
  - Edge Detection?
  - Texture Grouping?
  - Hough Transform?
    - Let's assume types of objects are unknown.




# Segmentation

Segmentation algorithms generally are based on one of two basis properties of intensity values.


- **Discontinuity:** to partition an image based on abrupt changes in intensity (such as edges).
- **Similarity:** to partition an image into regions that are similar according to a set of predefined criteria.




Parallelism



Continuity




Symmetry




Closure

# Segmentation needs Grouping/Clustering

- Grouping (or clustering)
  - Collect pixels that “belong together”







# Boundaries or Regions are equally good







# Boundaries or Regions are equally good





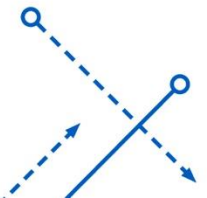
# Occlusion cues are important

- What do you see?



# Occlusion cues are important

- Aren't they?



# Regions and Edges

- Regions are bounded by closed contours
  - We could “fill” closed contours to obtain regions
  - We could “trace” regions to obtain edges
- Unfortunately, these procedures rarely produce satisfactory results.



# Regions and Edges

---

- **Edges** are found based on **DIFFERENCES** between values of adjacent pixels.
- **Regions** are found based on **SIMILARITIES** between values of adjacent pixels.
- We want group higher level units shared within regions of the image.

# Machine Learning

|                   | <i>Supervised Learning</i>       | <i>Unsupervised Learning</i> |
|-------------------|----------------------------------|------------------------------|
| <i>Discrete</i>   | classification or categorization | <div>clustering</div>        |
| <i>Continuous</i> | regression                       | dimensionality reduction     |

# Clustering example: image segmentation

- Break up the image into **meaningful** or **perceptually similar** regions





# Clustering example: image segmentation


- Foreground/Background









# Types of segmentations

Oversegmentation



Undersegmentation




Multiple Segmentations



# Major processes for segmentation

- Bottom-up
  - Group pixels with similar features
- Top-down
  - Group pixels that likely belong to the same object





Input



Bottom-up



Top-down



# Clustering

---

- Group together similar points and represent them with a single category.
- Key Challenges:
  - What makes two points/images/patches similar?
  - How do we compute an overall grouping from pairwise similarities?

# How do we cluster?

---

- K-means
  - Iteratively re-assign points to nearest cluster center
- Single-link clustering
  - Start with each point as its own cluster and iteratively merge the closest clusters
- Mean-shift clustering
  - Estimate modes of PDF

# Clustering for Summarization

- **A cluster:** a “center” (in feature space) and a list of data points it contains.
- Goal:
  - Minimize variance in data given clusters
  - Preserve information


$$\mathbf{c}^*, \delta^* = \operatorname{argmin}_{\mathbf{c}, \delta} \frac{1}{N} \sum_j^N \sum_i^K \delta_{ij} (\mathbf{c}_i - \mathbf{x}_j)^2$$

Cluster Center      Data

Whether  $\mathbf{x}_j$  is assigned to  $\mathbf{c}_i$

# K-means algorithm


1. Randomly select K centers (means)



2. Assign each point to nearest center (mean)




3. Compute new center (mean) for each cluster



# K-means algorithm


1. Randomly select K centers



2. Assign each point to nearest center



3. Compute new center (mean) for each cluster




Back to 2





# K-means algorithm



# K-means algorithm

---

1. Initialize cluster centers:  $\mathbf{c}^0$  ;  $t=0$

2. Assign each point to the closest center

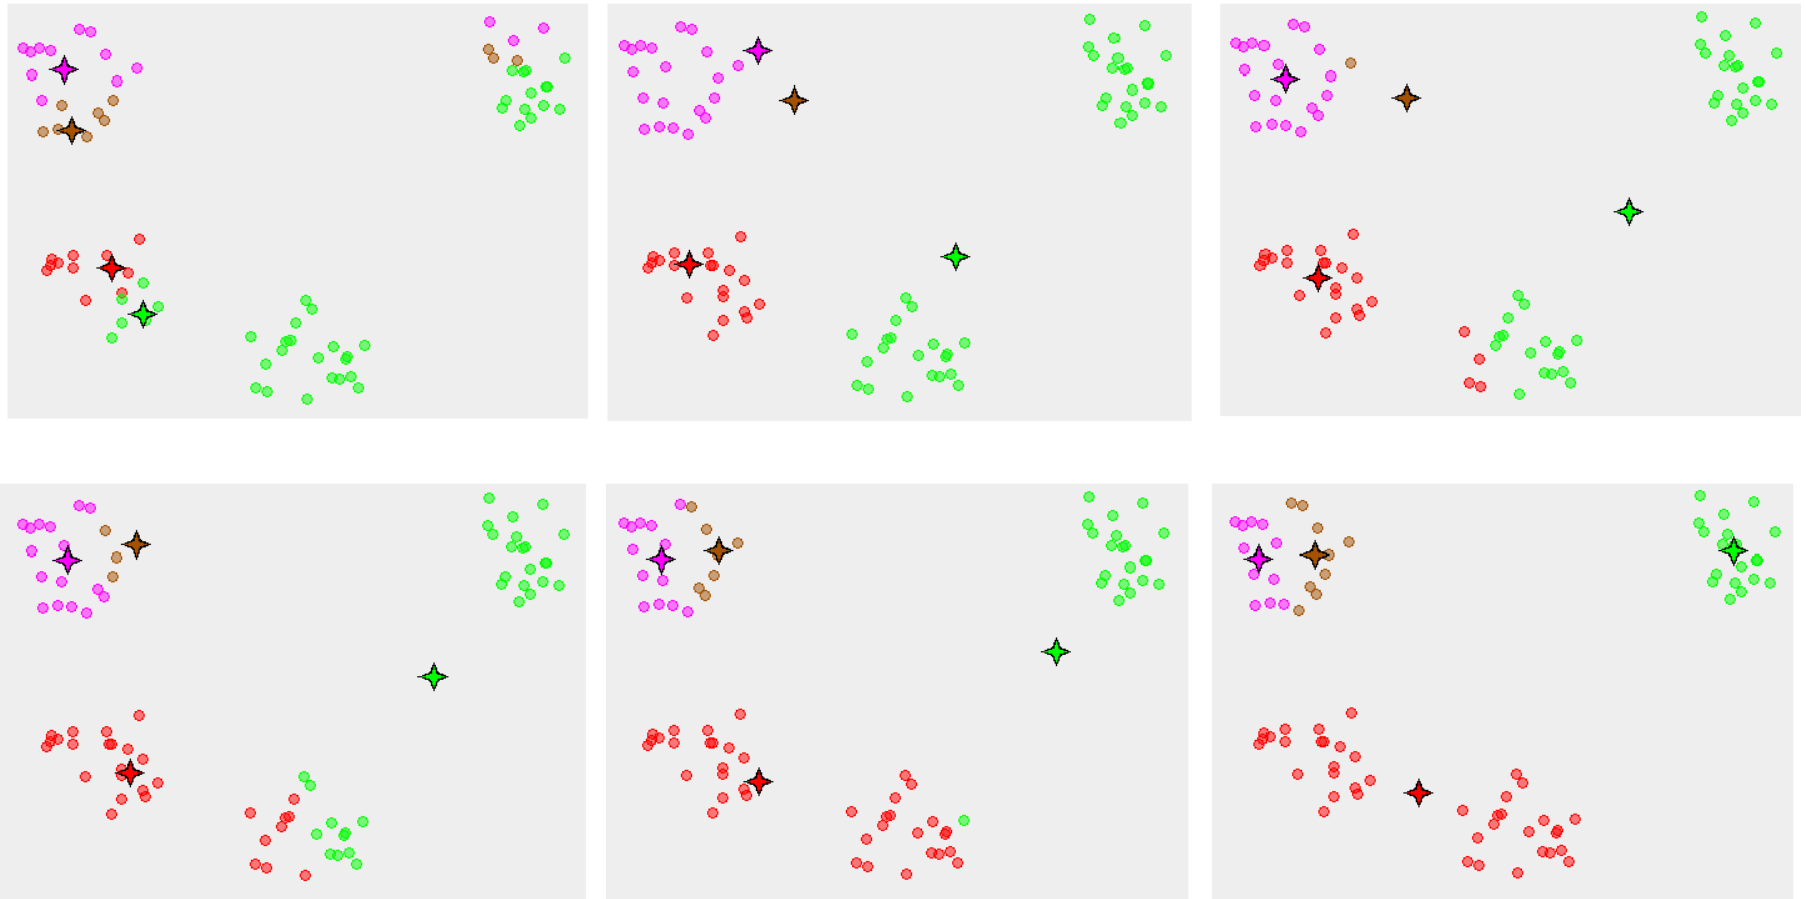
$$\delta^t = \underset{\delta}{\operatorname{argmin}} \frac{1}{N} \sum_j \sum_i^K \delta_{ij} \left( \mathbf{c}_i^{t-1} - \mathbf{x}_j \right)^2$$

3. Update cluster centers as the mean of the points

$$\mathbf{c}^t = \underset{\mathbf{c}}{\operatorname{argmin}} \frac{1}{N} \sum_j \sum_i^K \delta_{ij}^t \left( \mathbf{c}_i - \mathbf{x}_j \right)^2$$

4. Repeat 2-3 until no points are re-assigned ( $t=t+1$ )

# K-means may converge to local minimum




# K-means: design choices

---

- Initialization
  - Randomly select  $K$  points as initial cluster center
  - Or greedily choose  $K$  points to minimize residual
- Distance measures
  - Traditionally Euclidean, could be others
- Optimization
  - May converge to a *local minimum*
  - May want to perform **multiple restarts**

# K-Means pros and cons

- Pros
  - Finds cluster centers that minimize conditional variance
    - good representation of data
  - Simple and fast\*
  - Easy to implement
- Cons
  - Need to choose K
  - Sensitive to outliers
  - Prone to local minima
  - All clusters have the same parameters (e.g., distance measure is non-adaptive)
  - \*Can be slow: each iteration is  $O(KNd)$  for N d-dimensional points
- Usage
  - Rarely used for pixel segmentation




(B): Ideal clusters



# How to choose the number of clusters?





---

- On validation set
- Try different numbers of clusters and look at performance
  - When building dictionaries, more clusters typically work better



# Building Visual Dictionaries

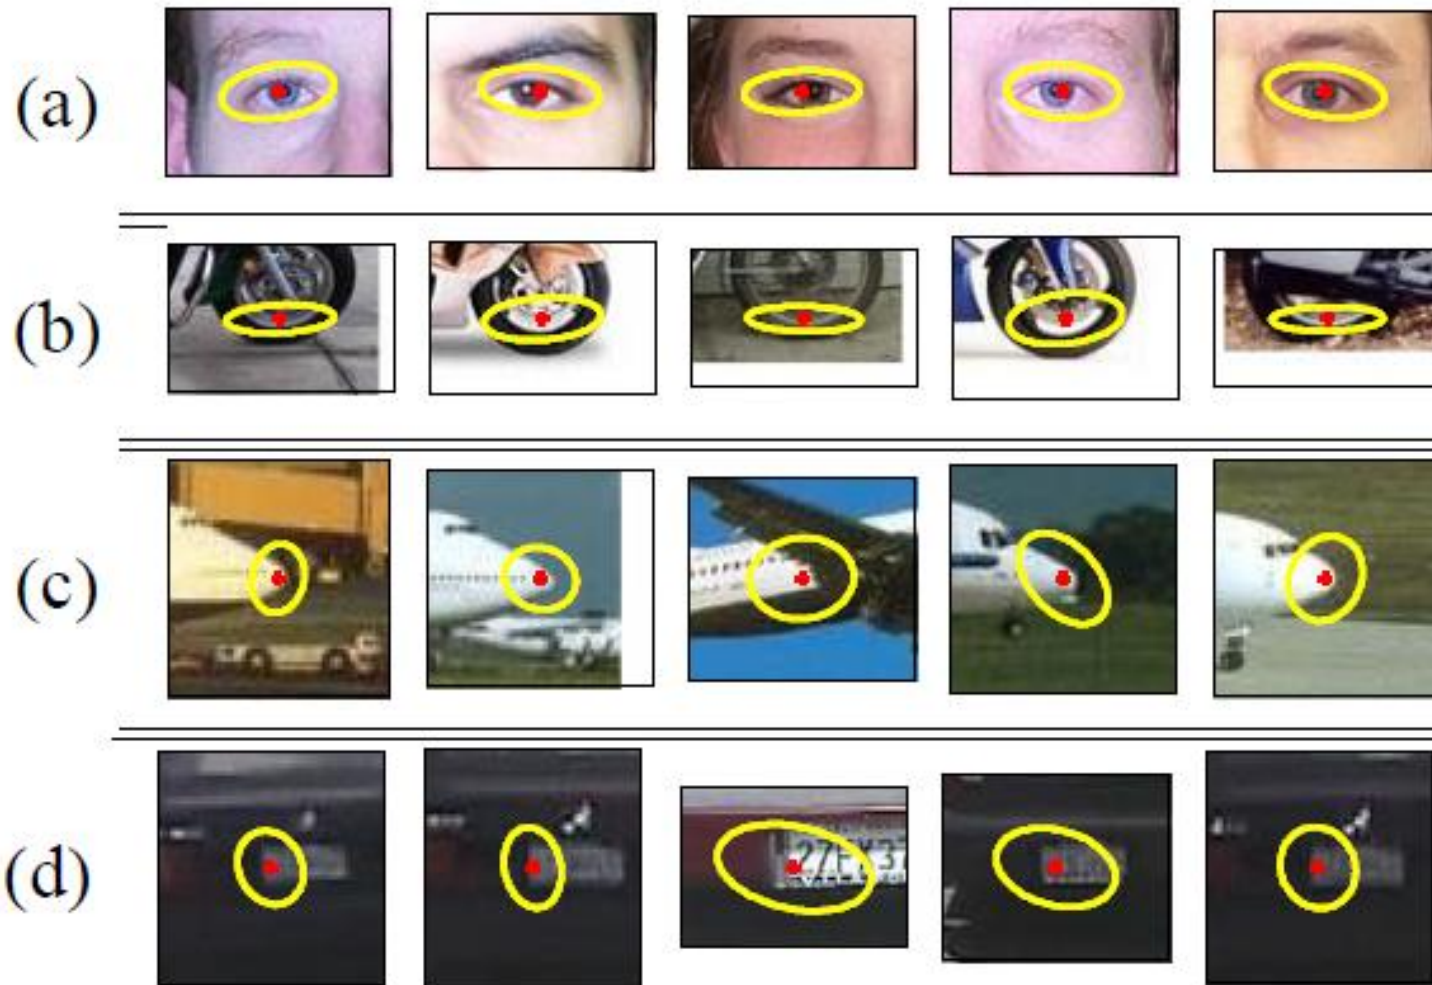
1. Sample patches from a database
  - E.g., 128 dim SIFT
2. Cluster the patches
  - Cluster centers are the dictionary
3. Assign a codeword (number) to each new patch, according to the nearest cluster



Codewords dictionary




# Examples of learned codewords




# How do we cluster?

---

- K-means
  - Iteratively re-assign points to the nearest cluster center
- Single-link clustering
  - Start with each point as its own cluster and iteratively merge the closest clusters
- Mean-shift clustering
  - Estimate modes of PDF
- Spectral clustering
  - Split the nodes in a graph-based on assigned links with similarity weights




# Single Link Clustering




1. Say "Every point is its own cluster"


# Single Link Clustering



1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters




# Single Link Clustering



1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster




# Single Link Clustering



1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster
4. Repeat



# Single Link Clustering




1. Say "Every point is its own cluster"
2. Find "most similar" pair of clusters
3. Merge it into a parent cluster
4. Repeat






# Single Link Clustering

- How many clusters?
- Threshold based on max number of clusters or based on distance between merges




# Single Link Clustering


- How to define cluster similarity?
  - Average/maximum/minimum distance of points
  - Distance between means.



$$L(r, s) = \max(D(x_{ri}, x_{sj}))$$



$$L(r, s) = \min(D(x_{ri}, x_{sj}))$$



$$L(r, s) = \frac{1}{n_r n_s} \sum_{i=1}^{n_r} \sum_{j=1}^{n_s} D(x_{ri}, x_{sj})$$

# Summary: Single Link Clustering

---

## Good

- Simple to implement, widespread application
- Clusters have adaptive shapes
- Provides a hierarchy of clusters


## Bad

- May have **imbalanced** clusters
- Still have to choose number of clusters or threshold
- Need to use an “ultrametric” to get a meaningful hierarchy

# How do we cluster?

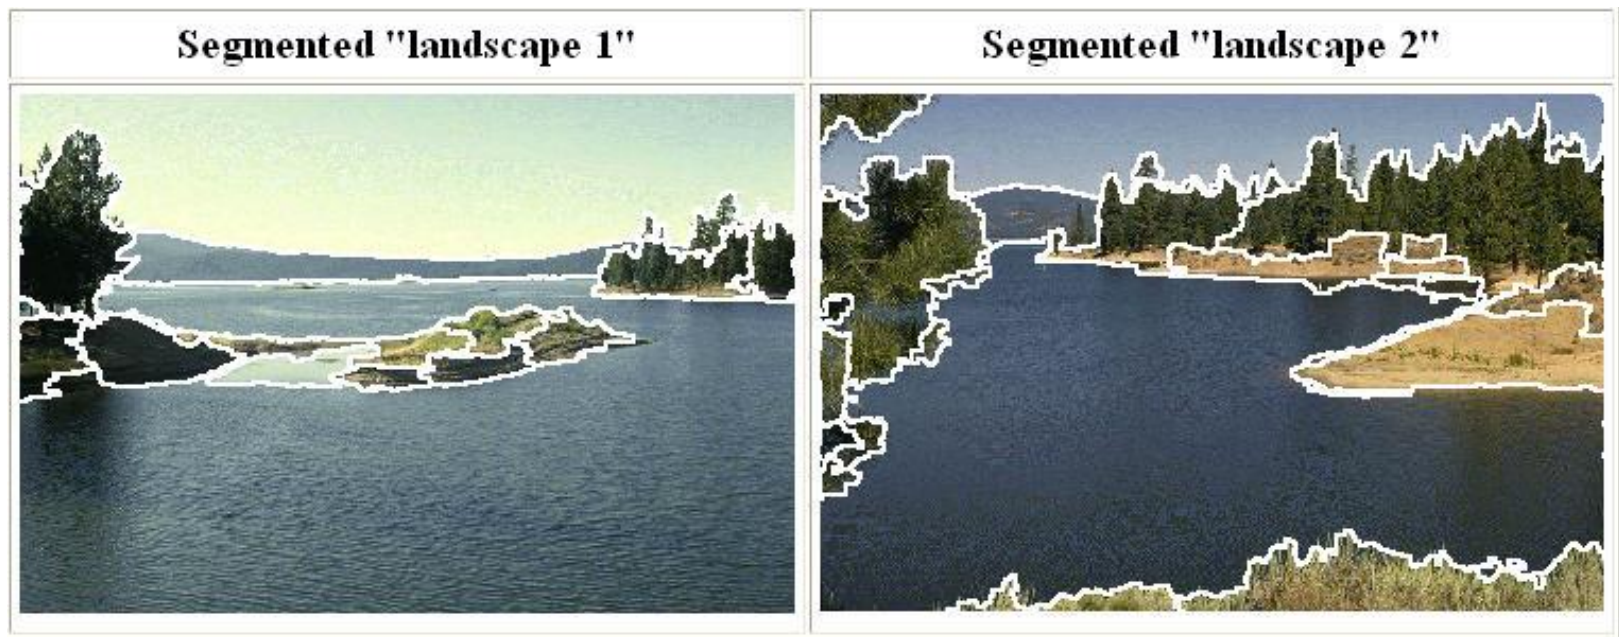
---

- K-means
  - Iteratively re-assign points to the nearest cluster center
- Single-link clustering
  - Start with each point as its own cluster and iteratively merge the closest clusters
- Mean-shift clustering
  - Estimate modes of PDF
- Spectral clustering
  - Split the nodes in a graph based on assigned links with similarity weights




# Mean shift algorithm

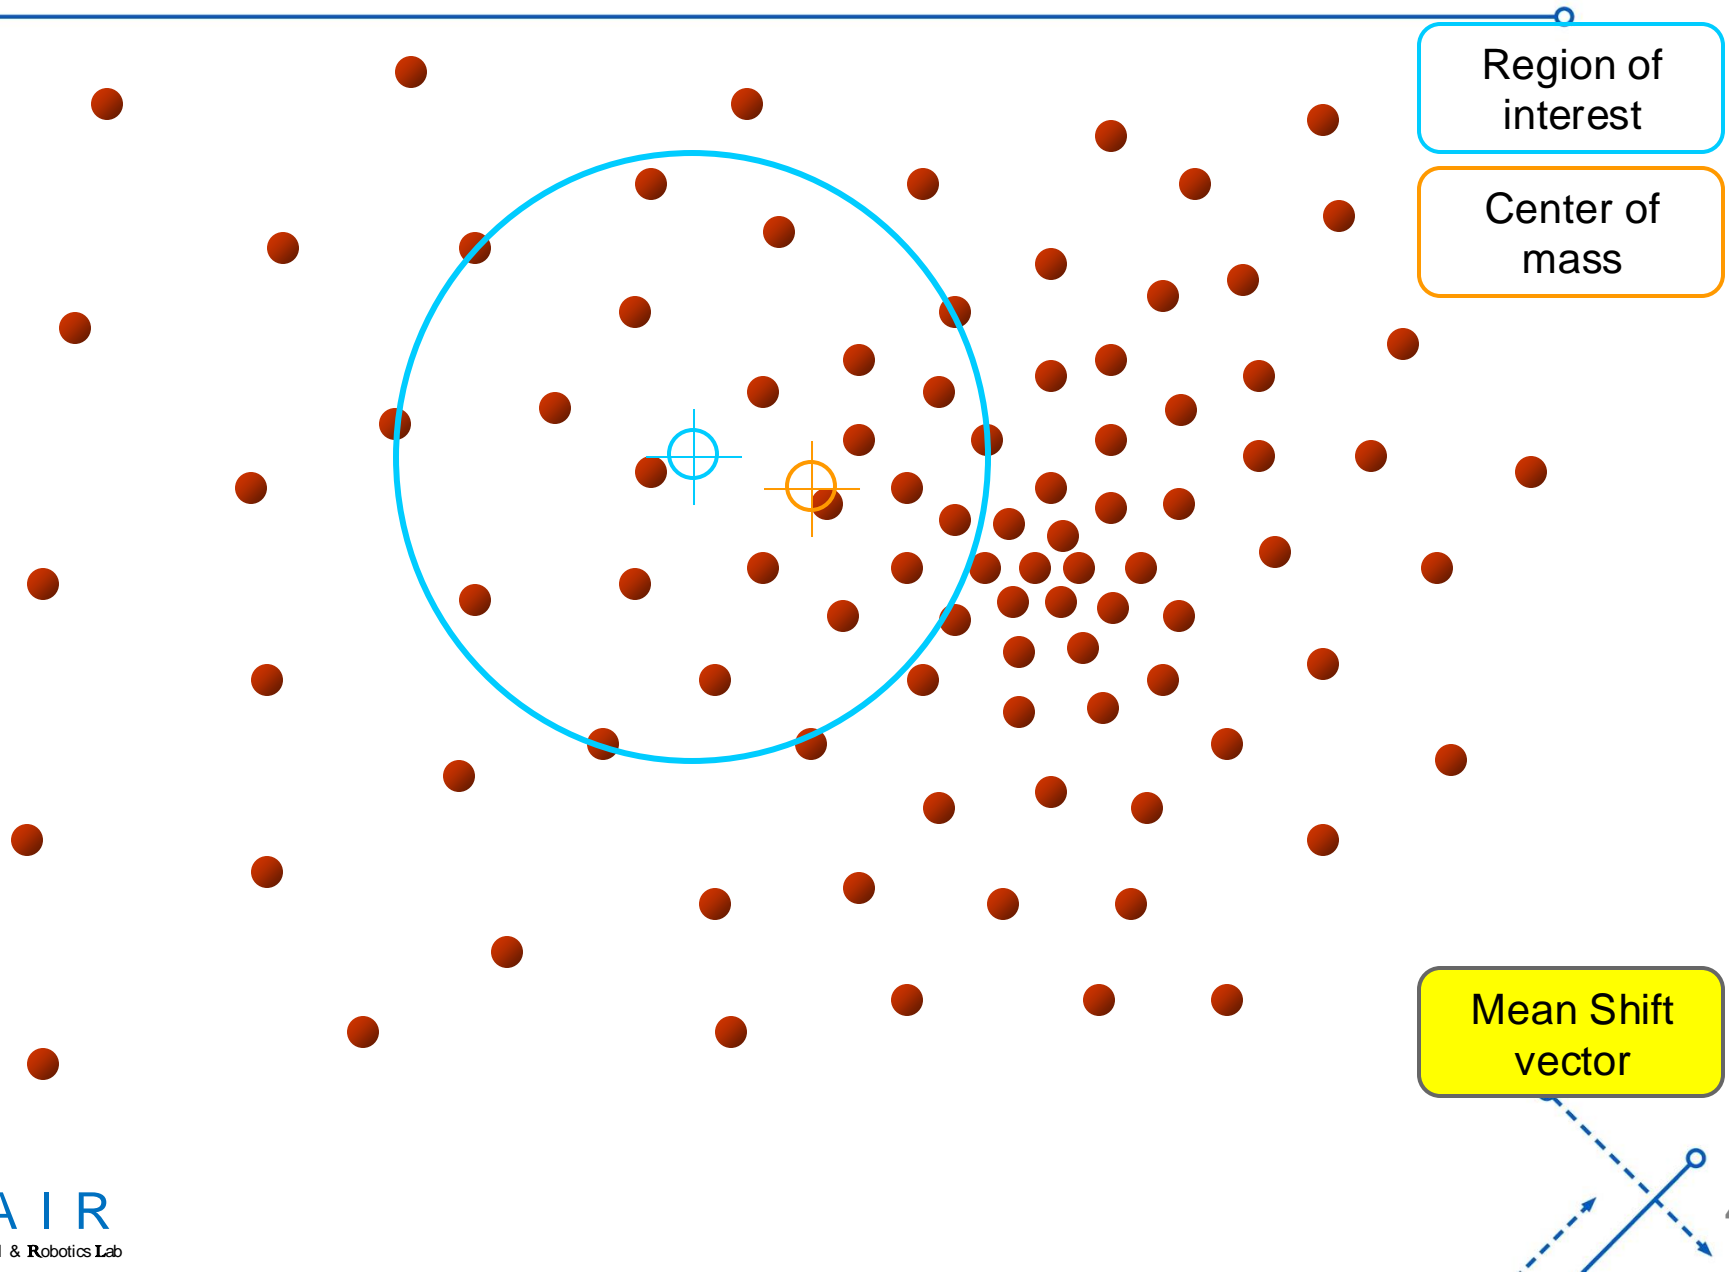
- Non-parametric feature-space analysis for locating the maxima of a density function.
- Versatile technique for clustering-based segmentation.




# Mean shift



# Mean shift




# Mean shift







# Mean shift




# Mean shift



# Mean shift

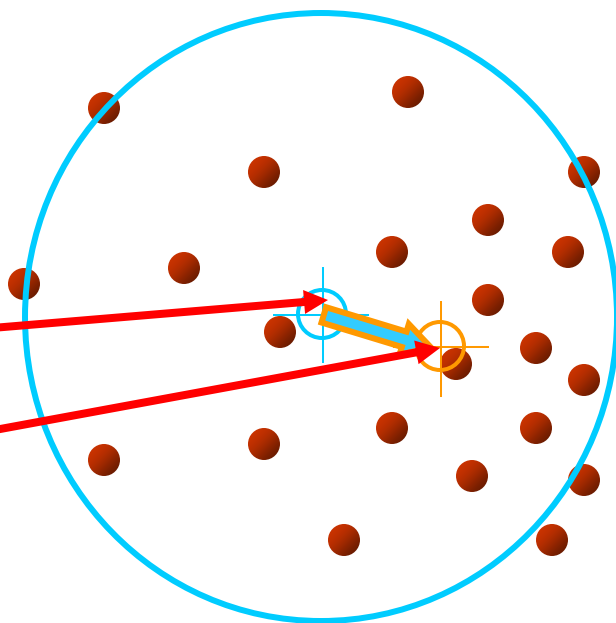


# Mean shift




# Computing the Mean Shift

- Compute mean shift vector.
- Translate mean by  $\mathbf{m}(\mathbf{x})$ , weighted by kernel function.

$$\mathbf{m}(\mathbf{x}) = \left[ \frac{\sum_{i=1}^n \mathbf{x}_i g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)}{\sum_{i=1}^n g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)} - \mathbf{x} \right]$$


$$g(x_i - x) = e^{-c\|x_i - x\|^2}$$

# Mean shift segmentation results



# Mean shift segmentation results





# Mean shift pros and cons

---

- Pros

- Good general-practice segmentation
- Flexible in number and shape of regions
- Robust to outliers

- Cons

- Have to choose kernel size in advance
  - Use kNN to determine window sizes adaptively
- Not suitable for high-dimensional features


- When to use it

- Oversegmentation, Multiple segmentations
- Tracking, clustering, filtering applications

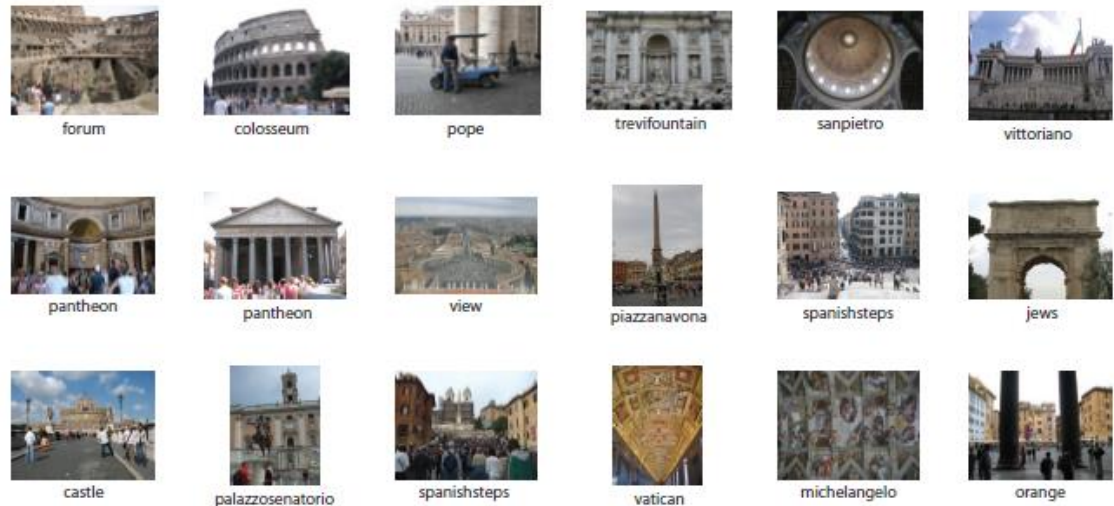


# Which algorithm to use?

- Quantization/Summarization: K-means
  - Aims to preserve variance of original data
  - Can easily assign new point to a cluster




Quantization for  
computing histograms



Summary of 20,000 photos of  
Rome using “greedy k-means”


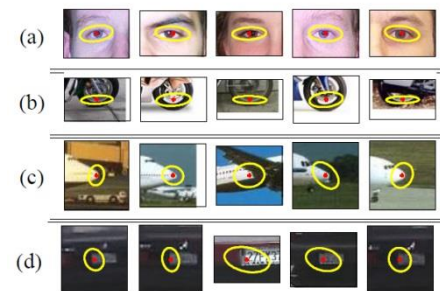
# Which algorithm to use?

- Image segmentation: Single Link
  - More flexible with distance measures
    - e.g., can be based on boundary prediction
  - Adapts better to specific data
  - Hierarchy can be useful



# Things to remember

- K-means useful for summarization, building dictionaries of patches, general clustering
- Single link clustering useful for segmentation, general clustering
- Mean-shift is useful for determining relevance, summarization, segmentation



# Application: Motion Segmentation

